

Product Requirements Document (PRD)

1. Executive Summary

Cosmiclly is a two-sided marketplace connecting **Clients** with **Advisors** (spiritual, professional, or coaching) for real-time 1:1 consultations. The platform supports **Text Chat**, **Audio Calls**, and **Video Calls**, facilitated by a secure "Pay-as-you-go" (Auth & Capture) payment system.

Core Value Proposition:

- **For Clients:** Instant, secure access to verified advisors via a unified communication platform. Users pay only for the exact minutes they use, with no requirement to hold pre-paid balances.
- **For Advisors:** A "Business-in-a-Box" to manage availability, set per-minute rates, and guarantee payment via Stripe Connect.
- **The Differentiator:** The "Twin AI" feature allows advisors to clone their knowledge and voice, providing 24/7 automated service to clients even when the human is offline.

2. User Roles & Permissions

ROLE	DESCRIPTION	KEY PERMISSIONS
Client	The end-user seeking advice.	Add payment method, browse advisors, initiate chat/calls, view session history, request data deletion.
Advisor	The professional providing service.	Create public profile, set per-minute rates, accept/decline requests, toggle Twin AI, view earning
Admin	Platform operator.	View global dashboards, process refunds, review advisor applications, manage featured content.

3. Functional Requirements

3.1 Authentication & Onboarding

- **FR-1.1:** Secure Signup/Login via Email & Password using Supabase Auth.
- **FR-1.2:** Role definition (**client**, **advisor**, **admin**) stored securely in user metadata or profiles table.
- **FR-1.3:** Advisor Application Flow: Prospective advisors must submit a bio, specialty, and pricing for Admin review before appearing in public searches.

3.2 Billing System (Auth & Capture)

- **FR-2.1: Stripe Connect Integration:** Advisors must link a bank account via Stripe Connect to receive payouts.
- **FR-2.2: The Pre-Auth Hold:** When a client initiates a call, the system prompts them to select a maximum duration (e.g., 15 mins). Stripe places a temporary hold on the client's card for the maximum potential cost (Duration × Advisor Rate).
- **FR-2.3: The Capture:** Upon call termination, the backend calculates the exact call duration and captures only the used amount, immediately releasing the remaining held funds.
- **FR-2.4: Insufficient Funds Prevention:** The LiveKit RTC token is strictly generated *only after* the Stripe hold is confirmed successful.

3.3 Marketplace Discovery

- **FR-3.1:** Browse advisors filtered by Category (Tarot, Astrology, etc.), Price, and Real-time Status (Online/Offline/In-Call).
- **FR-3.2:** View detailed Advisor Profiles including bios, specialties, reviews, and dynamic pricing displays.
- **FR-3.3:** Daily Horoscope integration providing automated, daily mystical content to drive recurring traffic.

3.4 Real-Time Communications (Human-to-Human)

- **FR-4.1: Unified Live Chat:** A persistent, WhatsApp-style chat thread per Client-Advisor pair utilizing Supabase Realtime, supporting typing indicators.
- **FR-4.2: Audio & Video Calls:** 1:1 secure WebRTC calls powered by LiveKit Cloud.
- **FR-4.3: Secure Signaling:** Access tokens are generated server-side via Supabase Edge Functions. Rooms are strictly isolated to the two authorized participants.
- **FR-4.4: Call State Machine:** Strict adherence to status transitions (**created** → **ringing** → **connected** → **ended**). Billing timers trigger precisely on the **connected** state.

3.5 The "Twin AI" (Digital Clone)

- **FR-5.1: Knowledge Base Ingestion:** Advisors can upload past readings and bios. The system vectorizes this text into a Supabase **pgvector** database.

- **FR-5.2: Text Mode (Offline Chatbot):** When an advisor is offline, the Twin AI automatically replies to incoming client chats using an LLM (e.g., OpenAI gpt-4o) grounded in the advisor's specific knowledge base.
- **FR-5.3: Voice Mode (Real-Time AI Call):** Clients can initiate a low-latency voice call with the AI Agent. The agent utilizes Cartesia for voice cloning to match the advisor's real voice, orchestrated via Vapi.ai.
- **FR-5.4: Tiered Pricing:** The platform supports separate per-minute rates for Human calls vs. Twin AI calls.

3.6 Admin Panel

- **FR-6.1:** Dedicated React UI protected by Row Level Security (RLS) for Admin roles.
- **FR-6.2:** Interface to approve or reject pending Advisor applications.
- **FR-6.3:** Ledger view to monitor active sessions, review flagged chats, and initiate Stripe refunds if disputes occur.

4. Non-Functional Requirements (Technical Constraints)

- **NFR-1: Data Privacy & RLS:** Clients can strictly only access their own data. Advisors can only access their assigned sessions.
- **NFR-2: Right to be Forgotten:** Automated API to anonymize chat data and delete PII upon user request to comply with GDPR/CCPA.
- **NFR-3: Performance:** LiveKit connection establishment must complete in under 2 seconds. Vapi voice-to-voice latency must target <800ms to ensure natural AI conversation flow.

DOCUMENT REVISION HISTORY

Version Number	Date	Revisions	Author
1.0	February 19, 2026	-Initial creation	Ivan Pamonag
1.1	February 26, 2026	-Changed ElevenLabs to Cartesia	Ivan Pamonag